

Lab 1

the solution of the assessment its in the bottom

Ali Jafar Altaweel 2240006456 7MA2

```
In [1]: # Step#3 + Step#4 : check that all the required packages are installed
import sys
import numpy as np
import scipy
import cv2
import skimage
import PIL
import matplotlib


print("Python      :", sys.version.split()[0])
print("NumPy      :", np.__version__)
print("SciPy       :", scipy.__version__)
print("OpenCV       :", cv2.__version__)
print("scikit-image:", skimage.__version__)
print("Pillow       :", PIL.__version__)
print("Matplotlib   :", matplotlib.__version__)
```

```
Python      : 3.13.1
NumPy       : 2.2.2
SciPy       : 1.16.3
OpenCV      : 5.0.0
scikit-image: 0.26.0
Pillow      : 12.3.0
Matplotlib  : 3.10.8
```

1. Loading and displaying an image with OpenCV

```
In [2]: import cv2
img = cv2.imread('./images/cameraman.tif') # Load the image data into a NumPy
array

import matplotlib.pyplot as plt
plt.imshow(img)
plt.show()
```



```
In [3]: print("Type of img :", type(img))
        print("img.shape  :", img.shape, " -> (rows, columns, channels)")
        print("img.dtype   :", img.dtype)
```

```
Type of img  : <class 'numpy.ndarray'>
img.shape    : (512, 512, 3)  -> (rows, columns, channels)
img.dtype    : uint8
```

```
In [4]: plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
        plt.title('cameraman.tif (BGR converted to RGB)')
        plt.axis('off')
        plt.show()
```


cameraman.tif (BGR converted to RGB)



2. Loading and displaying an image with PIL

```
In [5]: from PIL import Image
img2 = Image.open('./images/lena_gray_256.tif') # Load image into a PIL Image object

import matplotlib.pyplot as plt
import matplotlib.cm as cm
plt.imshow(img2, cmap = cm.Greys_r)
plt.show()
```



```
In [6]: print("Type of img2 :", type(img2))
print("img2.size      :", img2.size, "  ->  (width, height)")
print("img2.mode      :", img2.mode, "  ->  'L' means 8-bit grayscale")
print("img2.format    :", img2.format)
```

```
Type of img2 : <class 'PIL.TiffImagePlugin.TiffImageFile'>
img2.size    : (256, 256)  ->  (width, height)
img2.mode     : L        ->  'L' means 8-bit grayscale
img2.format   : TIFF
```

3. Saving an image to disk

```
In [7]: cv2.imwrite('new_image.jpg', img)
```

```
Out[7]: True
```

```
In [8]: img2.save('new_image2.jpg')
```

```
In [9]: # check that both files were really created
import os
for f in ['new_image.jpg', 'new_image2.jpg']:
    print(f, "->", os.path.getsize(f), "bytes", " exists:", os.path.exists(f))
```

```
new_image.jpg -> 68680 bytes  exists: True
new_image2.jpg -> 11433 bytes  exists: True
```

4. Displaying the image as an array

```
In [10]: print(img.shape)  # Note: img loaded using OpenCV is a numpy array
print(img)
```

```

(512, 512, 3)
[[[156 156 156]
  [157 157 157]
  [160 160 160]
  ...
  [152 152 152]
  [152 152 152]
  [152 152 152]]

 [[156 156 156]
  [157 157 157]
  [159 159 159]
  ...
  [152 152 152]
  [152 152 152]
  [152 152 152]]

 [[158 158 158]
  [157 157 157]
  [156 156 156]
  ...
  [152 152 152]
  [152 152 152]
  [152 152 152]]

 ...

 [[121 121 121]
  [123 123 123]
  [126 126 126]
  ...
  [121 121 121]
  [113 113 113]
  [111 111 111]]

 [[121 121 121]
  [123 123 123]
  [126 126 126]
  ...
  [121 121 121]
  [113 113 113]
  [111 111 111]]

 [[121 121 121]
  [123 123 123]
  [126 126 126]
  ...
  [121 121 121]
  [113 113 113]
  [111 111 111]]]

```

```

In [11]: import numpy as np
          # Convert the image to a NumPy array
          img_array = np.array(img2)

          # Print the array
          print(img_array.shape)
          print(img_array)

```

```
(256, 256)
[[162 162 162 ... 168 171 155]
 [162 162 162 ... 168 171 155]
 [162 162 162 ... 168 171 155]
 ...
 [ 48  57  54 ...  75  97  88]
 [ 46  50  52 ...  89 100  98]
 [ 44  55  54 ...  97 100 105]]
```

5. Student tasks

Task 1

The student will be asked to perform different operations on the array using the NumPy library.

Task 1 Solution – NumPy operations on the image array

```
In [12]: import cv2
import numpy as np
import matplotlib.pyplot as plt

a = cv2.imread('./images/cameraman.tif', cv2.IMREAD_GRAYSCALE)  # single
channel this time

# ---- information about the array ----
print("shape :", a.shape, " size :", a.size, " ndim :", a.ndim, " dtype :",
a.dtype)
print("min =", a.min(), " max =", a.max(), " mean =", round(float(a.mean()),
2),
      " std =", round(float(a.std()), 2), " median =", np.median(a))
print("distinct gray levels :", len(np.unique(a)))
print("location of the darkest pixel :", tuple(int(v) for v in
np.unravel_index(a.argmin(), a.shape)))
```

```
shape : (512, 512)  size : 262144  ndim : 2  dtype : uint8
min = 0  max = 255  mean = 118.31  std = 62.0  median = 142.0
distinct gray levels : 256
location of the darkest pixel : (76, 221)
```

```
In [13]: # ---- indexing and slicing ----
print("pixel [100, 200] =", a[100, 200])
print("top-left 4x4 block :\n", a[:4, :4])

crop = a[60:260, 160:360]          # crop = a rectangular slice
sub = a[:, :4, :, :4]              # keep every 4th row and column
print("crop :", crop.shape, " sub-sampled :", sub.shape)
```


```
pixel [100, 200] = 9
top-left 4x4 block :
[[156 157 160 159]
 [156 157 159 158]
 [158 157 156 156]
 [160 157 154 154]]
crop : (200, 200)  sub-sampled : (128, 128)
```

```
In [14]: # ---- geometric operations are just array re-ordering ----
flip_v = a[::-1, :]          # same as np.flipud(a)
flip_h = a[:, ::-1]          # same as np.fliplr(a)
rot90 = np.rot90(a)
transposed = a.T

print("a[::-1] == np.flipud(a) :", np.array_equal(flip_v, np.flipud(a)))
print("rot90 four times gives back a :", np.array_equal(np.rot90(a, 4), a))

plt.figure(figsize=(16, 4))
for k, (im, t) in enumerate([(a, 'original'), (crop, 'crop a[60:260, 160:360]'),
                             (sub, 'a[:, :4, ::4]'), (flip_v, 'a[::-1, :] (flip up-down)'),
                             (flip_h, 'a[:, ::-1] (mirror)'), (rot90,
                             'np.rot90(a)')]):
    plt.subplot(1, 6, k + 1)
    plt.imshow(im, cmap='gray', vmin=0, vmax=255)
    plt.title(t, fontsize=9)
    plt.axis('off')
plt.tight_layout()
plt.show()
```

a[::-1] == np.flipud(a) : True
rot90 four times gives back a : True




```
In [15]: # ---- arithmetic and logical operations ----
negative = 255 - a
brighter = np.clip(a.astype(np.int16) + 80, 0, 255).astype(np.uint8) # clip
to avoid wrap-around
wrapped = a + np.uint8(80)                                           # uint8
wraps at 256
mask = a > 100                                                       #
boolean array
binary = np.where(mask, 255, 0).astype(np.uint8)
side_by_side = np.hstack([a, negative])

print("mean original :", round(float(a.mean()), 1), " mean negative :",
      round(float(negative.mean()), 1))
print("pixels brighter than 100 :", int(mask.sum()), f"({100 *
mask.mean():.1f}%)")
print("pixels that wrapped after + 80 :", int((a.astype(int) + 80 > 255).sum()))
print("hstack shape :", side_by_side.shape)

plt.figure(figsize=(16, 4))
for k, (im, t) in enumerate([(negative, '255 - a (negative)'), (brighter,
'clip(a + 80)'),
                             (wrapped, 'a + 80 (uint8 wraps)'), (binary,
                             'np.where(a > 100, 255, 0)')]):
    plt.subplot(1, 4, k + 1)
    plt.imshow(im, cmap='gray', vmin=0, vmax=255)
    plt.title(t, fontsize=10)
    plt.axis('off')
```

```
plt.tight_layout()
plt.show()
```

mean original : 118.3 mean negative : 136.7
 pixels brighter than 100 : 187322 (71.5%)
 pixels that wrapped after + 80 : 34304
 hstack shape : (512, 1024)




```
In [16]: # ---- reshaping and statistics along axes ----
flat = a.ravel()
print("flattened :", flat.shape, " reshaped back equal :",
      np.array_equal(flat.reshape(a.shape), a))

row_means = a.mean(axis=1)      # one value per row
col_means = a.mean(axis=0)      # one value per column
counts = np.bincount(flat, minlength=256)
print("brightest row :", row_means.argmax(), " darkest column :",
      col_means.argmin())
print("most frequent gray level :", counts.argmax(), "with", counts.max(),
      "pixels")

plt.figure(figsize=(15, 3.5))
plt.subplot(1, 3, 1); plt.plot(row_means); plt.title('mean of each row
(axis=1)'); plt.xlabel('row')
plt.subplot(1, 3, 2); plt.plot(col_means); plt.title('mean of each column
(axis=0)'); plt.xlabel('column')
plt.subplot(1, 3, 3); plt.bar(np.arange(256), counts, width=1);
plt.title('histogram with np.bincount'); plt.xlabel('gray level')
plt.tight_layout()
plt.show()
```

flattened : (262144,) reshaped back equal : True
 brightest row : 73 darkest column : 151
 most frequent gray level : 13 with 6212 pixels



Explanation: because the image is a NumPy array, every operation here is a single NumPy expression. Slicing crops and sub-samples, reversing an axis flips the image, arithmetic changes the brightness (as long as the `uint8` wrap-around is avoided with `np.clip`), and a comparison gives a boolean mask that is directly a thresholded image.